

BI NORWEGIAN BUSINESS SCHOOL

GRA 19702 PRELIMINARY MASTER THESIS REPORT

Predicting Small Cap Stock Returns

Date of submission:

June 7, 2021

University:

BI Oslo

Study Program:

Master of Science in Business Analytics

Supervisor:

Eliane Bucher

This paper is a part of the MSc programme at BI Norwegian Business School.

The school takes no responsibility for the methods used, results found and conclusions drawn.

Abstract

This report presents the tentative plan for assessing the topic of predicting small cap stock returns on the Oslo Stock Exchange. The aim of the paper is to explore the predictive performance of deep learning algorithms when applied to stock returns, and compare returns for both large cap stocks and small cap stocks. It presents previous research and literature within the field, and theory related to deep learning and neural networks. It also provides current suggestions on choice of research design and methodology, which includes the suggested algorithms to be implemented, how stock returns will be predicted and measured, and which variables to be incorporated in the model training and testing. Lastly, a tentative plan for data collection and analysis is presented.

Contents

| | Page |
|---|-------------|
| 1 Introduction | 1 |
| 1.1 Problem Definition | 3 |
| 1.2 Literature Review | 4 |
| 2 Theory | 7 |
| 2.1 The Efficient Market Hypothesis | 7 |
| 2.2 Mean Reversion | 10 |
| 2.3 Technical Analysis | 11 |
| 2.4 Machine Learning | 15 |
| 2.4.1 History | 15 |
| 2.5 Deep Learning | 18 |
| 2.5.1 Biological Neural Networks | 18 |
| 2.5.2 Neural Network Structure | 19 |
| 2.5.3 The Learning Process | 23 |
| 2.5.4 LSTM Network | 26 |
| 2.5.5 GRU | 29 |
| 2.5.6 CNN | 30 |
| 2.6 Backtesting | 32 |
| 2.7 Transaction Costs | 33 |
| 2.8 Market Capitalization | 35 |
| 2.8.1 Large Cap Stocks | 35 |
| 2.8.2 Small Cap Stocks | 36 |
| 3 Data | 37 |
| 4 Methodology | 38 |
| 4.1 Long Short-Term Memory Network | 39 |
| 4.1.1 Training and Testing | 39 |
| 4.1.2 Independent Variables | 39 |
| 4.1.3 Target Variable | 41 |
| 4.1.4 Hyperparameters | 41 |
| 4.2 Gated Recurrent Unit Network | 41 |

| | | |
|-------------------|--|-----------|
| 4.2.1 | Training and Testing | 41 |
| 4.2.2 | Independent Variables | 41 |
| 4.2.3 | Target Variable | 41 |
| 4.2.4 | Hyperparameters | 41 |
| 4.3 | Convolutional Neural Network | 41 |
| 4.3.1 | Training and Testing | 41 |
| 4.3.2 | Independent Variables | 41 |
| 4.3.3 | Target Variable | 41 |
| 4.3.4 | Hyperparameters | 41 |
| 4.4 | Performance Metrics | 41 |
| 4.5 | Strategies | 41 |
| 4.6 | Predictive Models | 41 |
| 4.6.1 | Benchmark Algorithms | 41 |
| 4.6.2 | Deep Learning Algorithms | 42 |
| References | | 43 |

List of Figures

| | | |
|-----|--|----|
| 2.1 | Illustration of a Neuron (source: Wikimedia Commons) | 19 |
| 2.2 | Illustration of a Neural Network | 21 |
| 2.3 | Weights in a Neural Network | 22 |
| 2.4 | Gradient Descent | 25 |
| 2.5 | Structure of a Simple RNN | 26 |
| 2.6 | Structure of a LSTM-Network (source: [Olah, 2015]) | 27 |
| 2.7 | The Repeating Module (source: [Olah, 2015]) | 28 |
| 2.8 | Structure of a Gated Recurrent unit (source: [Olah, 2015]) . . . | 29 |
| 2.9 | Structure of a CNN (source: [Zakka et al., 2021]) | 31 |

1. Introduction

Artificial intelligence, and specifically machine learning, has experienced rapid growth during the last decades. Several industries and companies are increasingly implementing the power of machine learning to assist their day-to-day business operations. The stock market is no exception, where hedge funds and professional traders are making use of computational power for algorithmic trading, stock price predictions and as support for decision-making. The concept of predicting future stock price movements has been a topic of discussion for many years. One view is that future stock prices, or at least future price movements, can be predicted through technical analysis, which relies on studying past stock prices to predict future movements, or through fundamental analysis, which consists of studying companies' financial information such as earnings, equity, liabilities, cash flow etc.

Critics argue that stock prices are unpredictable, where price movements follow a "random walk". The term is commonly used in financial literature to describe price series where changes in price represent random departures from previous prices [Malkiel, 2003]. The "random walk" idea is closely related to the efficient market hypothesis. The hypothesis is based on a belief that markets are completely efficient, meaning that all information and news related to individual stocks and the market as a whole are reflected in the prices without delay. An immediate stock price reflection means that future price movements can only be explained by future news flows and cannot be predicted based on historical prices nor past events. As a result of this, investors are not able to obtain a greater return than a randomly selected portfolio consisting of individual stocks and with a comparable level of risk. This also means that investors with little market experience are able to achieve the same rate of return as professional investors through holding a diversified portfolio with prices given by the market [Malkiel, 2003].

The main concepts within the efficient market hypothesis has been challenged by both technical - and fundamental analysis, where the belief is that future

stock prices and price directions can be predicted. During the last years the hypothesis has been further challenged by an increasing use of artificial intelligence and machine learning as a tool for predicting stock prices. There are several different machine learning techniques that can be applied to this area of study. A common approach is called deep learning, which differs from more traditional machine learning techniques, in the form of not only being able to improve decisions based on previous learning, but also make intelligent decisions without human interference. The basis of creating a machine learning algorithm consists of splitting historical data into training sets and test sets, where the model is repeatedly trained and tested to improve accuracy and performance. The aim is for the model to be able to find underlying relationships within the data in order for the model to make predictions on new and unseen data. Artificial neural networks (ANN) is a type of deep learning which tries to resemble how the human brain functions. Stock price predicting is a complex topic, which requires intricate algorithms in order to detect underlying patterns in the data. ANNs are therefore more suitable for stock price predicting than more simple machine learning models.

There exists currently some research on stock price prediction for companies that are categorized as large cap stocks. These types of companies are normally mature companies associated with high stable earnings, and with a low level of risk with regards to the company's ability to generate future cash flow. However, investing in large cap stocks does not necessarily lead to greater returns as they often have low volatility which therefore leads to minor changes in stock price. On the other hand, companies that are categorized as small cap stocks are often associated with high volatility and large fluctuations in stock price. Research on how small cap stocks perform compared to large cap stocks dates all the way back to 1981, where Banz found evidence of a trend referred to as the "size effect" [Banz, 1981]. He examined common stocks listed on NYSE in the period between 1936-1975 and found that smaller firms had, on average, higher risk adjusted returns than larger firms. The evidence has been further confirmed with similar research carried out by Bauman et al. in 1998, in which they examined data from 21 international markets in the period between 1986-1996 and found the same "size effect" for markets outside of the US [Bauman et al., 1998]. In a more recent research article from 2019 Rompotis, Gerasimos explores the same topic of interest with data on 66 large-cap and 34 small-cap USA-listed ETFs (Exchange Traded Funds) during the period 2012-2016. The findings support previous evidence of the "size effect", but a higher level of risk is highlighted as

a factor related to small cap ETFs [Rompotis, 2019].

The main contribution of the following paper consists of examining if it is possible to apply the same logic and techniques used for predicting stock returns for large cap companies to predict stock returns for small cap companies. The paper aims to advance current literature within the field, where the research scope is extended to include both large cap stocks and small cap stocks. The paper will incorporate several deep learning algorithms for predicting stock returns, and it will present findings on how well the models perform on both groups of stocks and see if small cap trading results in greater returns than large cap trading. Given that previous research suggests a higher return associated with holding a portfolio consisting of small cap stocks, the following research can potentially advance the field of stock predictions, reduce the added risk in small cap trading, and increase returns from stock investments.

1.1 Problem Definition

The recent development within artificial intelligence and deep learning, as well as the increasing use within several sectors and industries, presents opportunities within the field of finance to further explore the applicability of prediction models on stock investments. The aim of the thesis consists of developing several deep learning algorithms where model performance is evaluated on both small cap- and large cap stocks. The scope of the paper is set to two groups of stocks listed on the Oslo Stock Exchange, which can be characterized as small cap- and large cap stocks. The paper will also discuss how the prediction models can be implemented in a real-life trading strategy, and explore the potential complications of small cap trading. The research will be of interest for both academics and institutions within the financial industry, and aims to answer the following research questions:

1. *How well do neural networks predict price movements for small cap stocks on the Oslo Stock Exchange, and do they outperform large cap stocks in terms of greater returns?*
2. *Is it feasible to implement small cap predictive models into a real-life trading strategy?*

The trading strategy the models will be tested on consists of long positions, which means that short strategies are not included in the scope of the paper. The reasoning behind this is that the possibility of shorting will greatly complicate the research and go beyond the limitations of this paper.

1.2 Literature Review

Applying deep learning and artificial neural networks to financial markets and trading is a relatively new topic of interest, and the current research is mainly focused on predicting large cap stock returns. To the best of my knowledge there does not exist any literature on predicting stock returns for small cap stocks listed on the Oslo Stock Exchange, where model performance and stock returns are compared to large cap stocks. The following section presents previous research on deep learning and stock return predictions, and emphasizes how the paper seeks to advance the current literature and in which areas.

Even though the topic of interest only has experienced increased focus during the last few years, research on predicting stock returns using deep learning dates back to 1988, where Halbert White examined neural network modelling for decoding nonlinear regularities in asset price movements [White, 1988]. White's article challenged the view of markets being efficient, based on the idea that when technology evolves (e.g. development within artificial intelligence), former common human beliefs can be questioned. The article assessed daily returns from IBM common stocks, where 1000 observations were used for training purposes, and 500 before- and after samples were used to evaluate the neural network. The network can be characterized as a "simple" network, where the resulting findings were not able to find evidence against the simple efficient market hypothesis due to overfitting issues. Despite not being able to disconfirm the simple efficient market hypothesis, White's research indicated a possibility of utilizing deep learning as a tool for exploiting potential inefficiencies in stock markets.

In 1993 Kryzanowski et al. contributed to the literature with an article on artificial neural networks, where they examined if the network was able to discriminate between stocks that provide positive returns and stocks that provide negative returns [Kryzanowski et al., 1993]. A neural network was trained to learn the relationships between a company's stock return one year forward,

financial data on the company and its industry, and macroeconomic variables. The neural network consisted of a pattern-recognition algorithm and can be termed as a classification model. The resulting findings showed that the neural network was able to correctly classify 72% of the positive/negative returns, meaning that the classification algorithm showed promising potential for stock picking. However, since the data consisted of a limited sample of smaller companies over a short time period, the results remained open for discussion.

In more recent work, Krauss et al. advanced the literature with two articles published in 2017 and 2018. In 2017, Krauss, Do and Huck published research which tried to bridge the gap between academic finance and the financial industry, where academic finance had started to focus on transparency and lower frequency data at the expense of performance, while the financial industry had started to focus their attention on performance and high frequency data at the expense of transparency [Krauss et al., 2017]. The authors deployed deep neural networks, gradient-boosted trees and random forests on all of the S&P 500 index constituents, with an aim of predicting the probability for each stock to outperform the general market. The findings showed that the neural networks did not outperform the other machine learning models, and also that the models performed particularly well in situations of market turmoil, in other words, when the markets are influenced by financial crisis, bubbles etc. Another interesting finding was that after the year of 2001 daily returns declined. The authors reasoned this with an increased usage of machine learning algorithms during this period.

In the following year, Krauss and Fischer further examined the use of deep learning algorithms for financial time series predictions [Fischer and Krauss, 2018]. The authors deployed long short-term memory (LSTM) for predicting out-of-sample directional movements for the constituents stocks of the S&P 500, using data from the period between 1992-2015. Long short-term memory is a type of artificial recurrent neural network (RNN) and differs from more standard feed forward neural networks in terms of having feedback connections. This will be further explained in the theory section. The research was an extension from the research carried out the year before, where the aim was to examine if a recurrent neural network could outperform more traditional machine learning algorithms, which consisted of memory free classification methods like random forest, deep neural network, and logistic regression classifier. The article presented evidence of the recurrent

neural network outperforming all of the other models with a sharpe ratio of 5.8, as well as having the highest model accuracy. Much like the previous research article market efficiency started to rise after 2010, which resulted in a decline in the network's profitability. The research of Krauss et al. points to a well-known theory consisting of mean-reversion, which means that asset prices and historical returns revert to the mean of the entire data set. The findings questions the ability of achieving greater returns (over a longer period of time) associated with applying prediction models on financial time series.

There has previously not been extensive research on how deep learning can be applied to stock return predictions on the Oslo Stock Exchange. However, in the last few years there have been a few articles on the subject. In 2016, Olden conducted research on whether or not it is possible to use machine learning algorithms to make a profitable stock trading scheme from stocks listed on the Oslo Stock Exchange [Olden, 2016]. He applied a stacked ensemble learning technique, which consisted of combining several machine learning algorithms into a single algorithm, and attempted to predict daily movements of the 22 stocks with the highest turnover on the Oslo Stock Exchange, using a total of 37 machine learning techniques. The research did not find concluding evidence of achieving a profitable scheme over a longer period of time, nor that stacked ensemble techniques outperforms other machine learning techniques. However, the top performing algorithms were able to outperform the Oslo Benchmark Index during the time period.

Two years later, Lund and Løvås conducted similar research by employing deep learning for stock return prediction on the Oslo Stock Exchange [Lund and Løvås, 2018]. Their paper was inspired by the work of Krauss and Fischer, where they wanted to add on the previous research on long short-term memory networks by testing similar methods and techniques on a less liquid stock exchange. As an additional contribution to the existing literature they examined the models' applicability to real-life trading, in terms of evaluating model performance when taking into account transaction costs, strategies related to reducing transaction costs, and the bid-ask spread. The resulting Sharpe ratio of applying the long short-term memory model on simple long trading strategies before transaction costs was 3.25, whereas the Oslo Stock Exchange Benchmark Index had a Sharpe ratio of 0.30. The long short-term memory model outperformed all of the other benchmark models when looking at model performance and accuracy. As in previous research, the model

experienced declining excess returns towards the last years of the period studied, which points to mean-reversion. It does however differ from previous studies, in terms of being able to have excess returns throughout the whole period. When including transaction costs, the findings show that excess returns are lost in the bid-ask spread, and the model is only able to produce a Sharpe ratio of 0.37.

2. Theory

The following chapter presents well-established theory on the efficient market hypothesis, before addressing deep learning and neural networks. Lastly, theory on trading strategies are presented, and how deep learning algorithms can be implemented.

2.1 The Efficient Market Hypothesis

The efficient market hypothesis is an economic theory introduced a generation ago, which was widely accepted among financial academics. A general view is that markets are extremely efficient in terms of reflecting news and new information in stock prices without delay. The efficient market hypothesis represents the idea of a "random walk", where in a price series, all subsequent price changes have random departures from previous prices. Information and news is believed to be immediately reflected in stock prices, which means that future stock price changes reflect only future news and is therefore independent of previous price changes. In other words, there is no correlation between tomorrow's price change and today's price change. New information regarding a stock is considered to be unpredictable, meaning that price changes are also unpredictable and random. Thus, stock price predictions based on either technical or fundamental analysis does not enable an investor to earn excess returns without accepting excess risk [Malkiel, 2003].

The hypothesis is commonly interpreted in three forms, which consist of a weak, semi-strong and strong form of efficiency. The weak form represents a market efficiency where prices reflect previous historical prices and returns. This implies that an investor is not able to use technical analysis as a method for

earning excess returns. Semi-strong form of efficiency consists of prices reflecting publicly available information, e.g. quarterly reports, stock splits, issue of shares, etc. This implies that an investor is not able to make use of fundamental analysis to earn excess returns. Finally, a strong form of efficiency represents prices where all information related to a stock, including information only available to company insiders, is reflected [Fama, 1970].

Jensen's Alpha (α) is a term within finance which was introduced in the late 1960's by Michael Jensen. The term is a post-hoc metric which evaluates a portfolio's ability to earn above-average returns adjusted for market risk. A positive Alpha means that the portfolio is able to generate a higher rate of return than the market as a whole, or some sort of benchmark index, while a negative Alpha translates to an underperforming portfolio compared to the market rate of return [Metts et al., 2008]. Opposed to the more general term Alpha, Jensen's Alpha takes into consideration the capital asset pricing model (CAPM) market theory and adjusts for risk in its calculations. The formula for Jensen's Alpha is as follows:

$$\text{Jensen's Alpha} = \text{Portfolio Rate of Return} - \text{Benchmark Portfolio Rate of Return}$$

$$\text{Benchmark Portfolio Rate of Return} = R_f + \beta (R_m - R_f)$$

Where,

$$R_f = \text{Risk Free Rate of Return}$$

$$\beta = \text{Beta}$$

$$R_m = \text{Expected Market Return}$$

The risk free rate of return is a representation of an investment with zero risk. A common way to find the correct risk free rate of return is to subtract the inflation rate from the yield of a treasury bond, which matches the investment duration. The U.S. treasury bill is an often used measure for a risk-free rate since the broad market considers there to be zero probability of the government defaulting on their obligations [Chen, 2020]. Beta is a measure of a portfolio's systematic risk, which represents the risk component that an investor is not able to avoid. It is used for measuring a stock's volatility and how much it fluctuates relatively to

the rest of the market [Kenton, 2021]. Beta can be found through the following calculation:

$$\text{Beta} = \frac{\text{Covariance}(R_e, R_m)}{\text{Variance}(R_m)}$$

Where,

R_e = The return on an individual stock

R_m = The return on the overall market

Covariance = How changes in a stock's returns are related to changes in the market's return

Variance = How far the market's data points spread out from their average value

From an investor's point of view, the key points of the efficient market hypothesis suggest that it is not possible to generate consistent alpha, and therefore no point of actively managing a portfolio. This means that a portfolio manager is not able to earn excess returns through the use of technical analysis nor fundamental analysis.

Previous research has presented evidence both for and against the efficient market hypothesis. In 1978, Michael C. Jensen argued that there was no other proposition in economics with more supportive, empirical evidence than the efficient market hypothesis [Jensen, 1978]. However, in the article he acknowledged that times were moving forwards with better data- and econometric analysis which uncovered anomalies and inconsistency with previous evidence supporting the hypothesis. Further, in 1980 Grossman and Stiglitz addressed the issue of prices fully reflecting all available information. They argued that prices only partially reflect available information, where those who obtain information are adequately compensated [Grossman and Stiglitz, 1980]. If prices reflected all information, then there would be no financial incentive to obtain information. This would affect the market efficiency in terms of there not being arbitrators in the market, who have a function of keeping the market efficient.

The strong form of the efficient market hypothesis was corrected by Fama in

1991, which was previously based on the assumption of zero costs associated with obtaining information and carrying out transactions. It was assumed that only if transaction costs and costs of obtaining information were equal to zero, investors would have the incentive to keep trading until prices reflected all available information. In reality these costs are positive (greater than zero), which does not correspond with the previous assumption. Fama redefined this part of the strong form of market efficiency as stock prices reflecting available information until the marginal usefulness of information is equal to the marginal cost [Lekovic, 2018]. A few years later, Fama addressed the issues with anomalies and referred to them as a result of chance. With an expected value of above-average return of zero, anomalies in the form of under- and overreactions are due to chance. Fama also believed anomalies to be a result of incorrect methodology [Lekovic, 2018].

Being that there does not exist conclusive evidence either for or against the efficient market hypothesis, the aim of the paper is to test if the weak-form of the efficient market hypothesis holds or not. The hypothesis will be implicitly tested in terms of examining if it is possible to achieve above-average returns using deep learning-models that are based on technical analysis, and also if the application of deep learning differs between small cap- and large cap stocks.

2.2 Mean Reversion

Mean reversion is a financial theory where supporters of the theory believe that a stock's historical returns and price volatility will eventually revert to the long-term mean, or the mean for the entire dataset [Chen, 2021]. Investors who base their investing strategy on this theory try to utilize the fluctuations in stock prices, with an expectation of the price eventually converging to the average price over time. Thus, if a stock price drops abnormally low, the theory suggests that it is probably a good time to buy, while if a stock rises abnormally high it is most likely time to sell, or perhaps even go short. The theory relates to the idea of buying low and selling high, which is a strategy that can be applied to most securities. However, there is no guarantee that a stock price will converge to the average price level. Several changes could happen which can affect the future outlook of a stock, both at a company level, market segment level, or the market as a whole. This applies to both a more positive and negative outlook for the company, which can be exemplified by a new breakthrough technology or a pending lawsuit.

James M. Poterba conducted research on the topic in 1988 with his article titled "Mean reversion in stock prices: Evidence and Implications" [Poterba and Summers, 1988]. In the article he states that "if market and fundamental values diverge, but beyond some range the differences are eliminated by speculative forces, then stock prices will revert to their mean". Hence, if the market value exceeds a threshold where the price reaches a level which can not be explained by fundamental values, the price will start converging to the mean. When stocks are incorrectly priced and eventually reverted to the mean, returns need to be negatively serial correlated at some frequency. The resulting findings from the research showed a positive serial correlation in stock returns over short periods and negative correlation over longer time periods. The statistical tests could not consistently reject the random-walk hypothesis at high significant levels, but at an aggregated level the evidence pointed to transitory components having a high explanatory effect on the variance in returns.

The findings related to mean reversion could potentially strengthen the assumption of deep learning algorithms being able to find patterns within historical data. Based on the theory, including securities' moving averages should enable neural networks to make predictions of profitable entry- and exit price levels.


2.3 Technical Analysis

As mentioned in the previous section, technical analysis is an investing technique that contradicts the weak form of the efficient market hypothesis. The investing strategy is a tool which tries to predict future price movements based on historical prices and other technical indicators like trends, volume, moving averages, support- and resistance levels, and momentum. While long-term investors focus more on time spent in the market, investors who use technical analysis are more focused on timing the market. The tool is used for timing profitable entry- and exit points, and it usually consists of more frequent trading.


Investtech is a Norwegian company which specializes in behavioral finance and technical analysis of stocks. They offer stock analysis based on advanced mathematical models and statistical optimization. The analyses are a result from more than 20 years of research on stocks that are listed on the Nordic

stock exchanges. Their research shows supporting evidence of the ability to achieve above-average returns using technical analysis. Some of the resulting findings on various technical indicators include:


- **Rising trend** - stocks that are in a rising trend channel yield annual excess returns of 7.8% in the medium long term.




- **Falling trend** - stocks that are in a falling trend channel yield annual below-average returns of 11.4% in the medium long term.



- **Trend break up** - stocks in rising trend channels which break up through the ceiling of the channel yield 14.7% in annual excess returns in the medium long term.




- **Trend break down** - stocks in falling trend channels which break down through the channel floor yield 27.9% in annual below-average returns in the medium long term.




- **Positive volume balance** - stocks with a positive volume balance yield annual excess returns of 4.7% in the medium long term.




- **Negative volume balance** - stocks with a negative volume balance yield annual below-average returns of 7.6% in the medium long term.



- **Strong positive momentum** - stocks with strong positive momentum yield 11.4% in annual excess returns in the medium long term.



- **Strong negative momentum** - stocks with strong negative momentum yield 13.3% in annual below-average returns in the medium long term.



[Linløkken, 2021]

There are numerous different indicators and measures within the field of technical analysis, and the importance of each one varies greatly between investors. However, the main object of this type of investing strategy is to capture both the market psychology as a whole and also for individual stocks. Investors who utilize technical analysis seek to identify levels of a stock price where there are likely more buyers than sellers (will often lead to an increase in price) and levels where there are likely more sellers than buyers (will often lead to a decrease in price). A simplified explanation of the directional movements in price consists therefore of stock prices being greatly dependent on supply and demand. The level of aggressiveness among buyers and sellers is another factor which affects the movements of a stock price, and captures the psychological state of mind of the investors. Sellers who close their position by aggressively selling at bid, represent a situation where there is little optimism among investors and a probable decline in short-term price. Buyers who open positions by aggressively buying at ask, represents a situation where there is a lot of optimism among the investors and anticipations of a short-term increase in the stock price.

Investors (traders) who base their investing strategy on technical analysis believe that the market psychology can be found in historical charts, prices and other technical measures, and that history repeats itself. An example of this is if a stock has previously reversed when falling to a certain price level on multiple occasions, the chances are that this will happen again. This is called a support level and represents a situation where there are many investors who believe the stock is cheap, and few sellers because they feel the stock is undervalued. Hence, there is a high probability that the stock price will increase based on a revitalized optimism among the investors again. This is under the assumption that there has not been any fundamental news that changes the company's future outlook.

In order for technical analysis to work as either an investing strategy or investing tool, there must exist patterns within historical data that can be taken advantage of. If investors, through the use of technical analysis, are able to achieve above-average returns, it means that there must exist certain inefficiencies and that the weak form of the efficient market hypothesis does not hold. Instead of manually having to search for and detect these inefficiencies, this paper aims to utilize

deep learning algorithms as an automated substitute for finding patterns within the data, representing potentially valuable entry points and ultimately above-average returns. Theory on deep learning is explained in the next section.

2.4 Machine Learning

Modern enterprises today utilize the power of machine learning and computational power to assist every day-operations and operational decisions. Implementing analytics as an automated decision tool throughout the entire organization was previously characterized by tech-companies like Google, Amazon, Facebook etc, but in recent years this trend has changed. More and more mainstream companies are trying to replicate the success of these tech-companies by developing predictive models, visualization models, and other types of data-driven models to support operational decisions. Even though analytics and machine learning just recently has started to be recognized as a business asset, and not just a separate business project, the concepts of machine learning have been around for many years.

2.4.1 History

Machine learning is partly based on the brain cell interaction-model created in 1949 by Donald Hebb in the book titled "The Organization of Behaviour". Hebb presented theories on neuron excitement and communication between neurons, which can be translated into the concepts of artificial neural networks and artificial neurons [Foote, 2019]. In the 1950's Arthur Samuel came up with the phrase "Machine Learning". He developed a chess computer program based on alpha-beta pruning, where the program chooses the next move by using a minimax strategy and a scoring function. Samuel introduced rote learning, where the program recorded all previous positions and combined this with the reward function. In the same decade, Frank Rosenblatt built upon the work of Hebb and Samuel by creating the perceptron. The program was designed for image recognition. However, the program did not produce satisfying results, where it was only able to recognize a few kinds of visual patterns. This led to a stagnation in the research of machine learning and neural networks, which lasted until the 1990's.

In 1967, Marcello Pelillo was given credit for developing the nearest neighbour algorithm. The algorithm was developed in connection with the traveling

salesperson's problem, where the algorithm was used for mapping routes and finding the most efficient one. This is considered as the beginning of basic pattern recognition. Further in the same decade, multilayers were introduced in the research of neural networks. Using more than one layer in the perceptron resulted in a significant increase in processing power, which laid the foundation for the future development of feedforward neural networks and backpropagation. The backpropagation algorithm differs from feedforward neural networks by having the ability to process output errors and distribute them backwards through the layers of the network. This improves the learning ability of the network and is currently being used to train deep neural networks [Foote, 2019].

During the period from the 1970's and 1990's the fields of machine learning and artificial intelligence were separated. Machine learning had previously been incorporated as a training program for artificial intelligence, but shifted focus to solve more service-oriented tasks during this period. Machine learning flourished in the 1990's as a result of the internet growth and focus on neural networks. Increased access to digital data and an easier ability to share data and services was the main reasoning behind the revived success of machine learning during this decade. In this decade, Robert Schapire introduced the concept of boosting. This is a technique developed to reduce bias during supervised learning and to transform weak learners into strong learners. Boosting algorithms create a final strong classifier by repetitive learning weak classifiers. Each weak classifier is weighted based on the learner's accuracy and then reweighted. The weight of misclassified input data increases, while the weight of correctly classified data decreases, in order for future weak learners to focus more on previous misclassified weak learners. Strong learners are easily classified aligned with the true classification, while weak learners are only slightly correlated with the true classification [Foote, 2019].

Ever since the millennium, machine learning has been applied to several business areas. In the early 2000's, speech and facial recognition technology was developed. Speech recognition is based on Long Short-Term Memory (LSTM), which is a neural network technique developed by Jürgen Schmidhuber and Sepp Hochreiter in 1997. LSTM is particularly relevant for speech recognition in terms of being able to learn tasks that require the model to be able to memorize discrete steps that took place thousands of steps earlier. The relevance of the model was highlighted in 2015 when Google's speech

recognition program gained a 49% increase in performance by using a LSTM model. During the same period, newly developed facial recognition algorithms were found to be 10 times more accurate than algorithms from 2002 and 100 times more accurate compared to algorithms applied in 1995 [Foote, 2019].

Today, machine learning is considered as the computer's ability of making automated actions without being explicitly programmed. The applicability of machine learning is relevant for several business areas, which include:

- **Self-driving vehicles** - Autonomous cars are currently being developed with the use of machine learning and deep learning. These cars have the ability to steer, accelerate and brake, based on prior knowledge, stimuli and past experiences. Prior knowledge includes street maps, the meaning of signs, what to stop for etc. Stimuli are made up of a combination of sensors like vision, lidar, radar, GPS and voice commands, while past experiences can for example be how braking and steering affects direction and speed. where prior knowledge like street maps, meaning of signs, what to stop for etc.
- **Business decisions** - Machine learning is increasingly being used for operational decisions, which covers various business operations like predicting production metrics, the number of new customers based on marketing campaigns, or predicting inventory based on customer demand.
- **Fraud detection** -Fraud related to tax evasion, money laundering and insurance are common problems for most countries, and cost governments billions of dollars each year. By implementing machine learning and predictive models, governments are now able to prevent and detect fraud in a much more efficient way.
- **Natural language processing** - Machine learning enables Email filtering, where algorithms are able to detect spam and recognize different categories the Email belongs to. Other examples of natural language processing include smart assistants, improvement in search results where search engines are able to give relevant results based on similar search behaviour, and predictive text like autocorrect and autocomplete.

Within finance, there is an increasing amount of hedge funds that are utilizing

machine learning to get an edge over the broad market. Hedge funds use it as support to anticipate market movements and detect market trends at an early stage. Machine learning models also help them get a better understanding of how specific news will affect the market direction and intensity. The main point by incorporating machine learning in this field is to be able to extract information out of large amounts of data. This can for example be to use satellite data to determine customer traffic in retail or to calculate the level of crude oil storage, or to use natural language processing to analyze market sentiment through large amounts of text related to important macroeconomic factors.

2.5 Deep Learning

In the field of artificial intelligence, deep learning is a subset of machine learning where neural networks can identify patterns within training data, where they eventually are able to make predictions from new, unseen data. The structure of deep learning is based on the workings of a human brain, where the brain is used to recognize patterns and classify various types of information. Neural networks can be trained to perform several tasks, which include regression, classification and clustering [Oppermann, 2019]. Before going more into detail on the network structure and how the network learns, it is important to get an idea of how the human brain functions, which in turn can give a better understanding of how deep learning networks work in practice.

2.5.1 Biological Neural Networks

Artificial neural networks can be compared to biological neural networks in terms of artificial neural networks trying to simulate some of the functionalities of the neural networks in the human brain. The processes are quite similar, where artificial neural networks try to imitate the human brain's biological neurons, but in a more simplified way.




Figure 2.1: Illustration of a Neuron (source: Wikimedia Commons)

The figure above illustrates the workings of a nerve cell (neuron) in the human brain. The neurons are arranged together to form a network of nerves, where they transmit and process information received from human senses. These are electrical impulses being passed from one neuron to another. Impulses from the synapse of an adjoining neuron are received through the part of the neuron termed dendrites. The impulses are then carried further to the nucleus (soma), where the electrical impulses are processed before being passed to the axon. The axon consists of a longer branch with a main function of carrying the electrical signal from the soma to the synapse. Lastly, the signal is passed to a neighboring neuron's dendrites, which creates an intricate network of neurons in the human brain [Panchal, 2018]. The neurons are electrically excitable due to their membranes' maintenance of voltage gradients. A large amount of change in voltage, in a short amount of time, creates an action potential (electrical pulse) that travels rapidly through the axon and activates adjoining neuron's synaptic connections [Oppermann, 2019].

2.5.2 Neural Network Structure

As with linear regression, the data consist of independent variables x_n , with an assumption that these features are able to explain the variability in the dependent variable y . The target variable can either be continuous or categorical, depending on whether the network is a regression or classification model. Categorical dependent variables consist of a finite number of categories or groups. Examples of a classification model could be predicting gender, type of animals, if an email is spam or not, or classifying customer churn based on recent behaviour. Classification models normally consists of four types of classification, which include:

- Binary classification - This is a very common classification task consisting of having two class labels, where one class represents the normal state and

the other class the abnormal state. The class representing the normal state is assigned the value 0, while the class representing the abnormal state is assigned the value 1. Binary classification models usually include a discrete probability distribution, where the models predict the probability of an observation belonging to the two classes.

- **Multi-class classification** - This type of classification refers to tasks where there are more than two class labels. This means that there could be thousands of known classes, as for example would be the case in face recognition. Normally, a multinoulli probability distribution is incorporated in the predictive model, where the model predicts the probability of an observation belonging to each of the several class labels.
- **Multi-label classification** - In this case, there are at least two class labels, where the classification task may require predicting more than one class label for each observation. This might be the case in photo classification, where there could potentially be several objects (class labels) in one photo (observation). Also in this type of classification task, it is common to incorporate a discrete probability distribution, where the model predicts multiple outputs and multiple binary classification predictions for each observation.
- **Imbalanced classification** - The last type of the most common classification problems is imbalanced classification. This type of classification refers to tasks where the number of observations in each class are unequally distributed and where most of the observations belong to the normal class label. A common technique to solve these types of problems consists of combining a normal binary classification model with either undersampling the majority class or oversampling the minority class.

[Brownlee, 2020]

Continuous dependent variables have infinite numeric values, which can also consist of date or time. Examples of regression (continuous output) could be predicting next month's sales, traveling time for a specific route, or the number of new customers based on a marketing campaign. In linear regression, coefficients for the independent variables are estimated to make predictions for the target variable, while neural networks incorporate non-linearity and a more intricate structure to gain a deeper understanding of patterns within the data.




Figure 2.2: Illustration of a Neural Network

Figure 2.2 illustrates a traditional neural network, where the nodes represent artificial neurons. These can be compared to the biological neurons of the human brain. The nodes are a graphical representation of numeric values, while the connections can be thought of as the axons in the biological neuron. These connections have different dynamic weights that change as the network trains to learn the patterns within the data. The strength of the connections also changes with an aim for the network to set the right weights, in order to make accurate predictions [Oppermann, 2019]. The layers of the neural network can be divided into an input layer, hidden layers and an output layer. The input layer is where the network receives the input data x , which consist of features considered to have an explanatory effect on the dependent variable. Each x is an entire vector where an input neuron represents one element in the vector. The hidden layers are where mathematical operations are performed in order for the neural network to obtain a prediction vector y . Lastly, the output layer is where the network's predictions are being made. An output neuron consists of the value of the prediction, which is a vector y . The values presented in this layer

could either be a probability value between 0 and 1 (classification) or an infinite value predicting a desired target value (regression) [Oppermann, 2019].




Figure 2.3: Weights in a Neural Network

Figure 2.3 shows an example of a simple neural network with an input layer consisting of two neurons and an output layer with 3 neurons. Each connection between two neurons is represented by a weight w , where each weight has indices. The first number of indices tells us the number of neurons from the connection's origin, while the second value tells us the number of neurons in the layer where the connection leads. The weights between the layers of a neural network can be illustrated through a weight matrix, which is shown below.

$$W = \begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{pmatrix}$$

The number of entries in a weight matrix corresponds to the number of connections between neurons. This means that the dimensions of the matrix is determined by the layers' size and their connections. The weight matrix above has two rows, since there are two neurons in the input layer and three columns because of the three neurons in the output layer [Oppermann, 2019].

2.5.3 The Learning Process

Having established the basic concepts and structures of a neural network, this section will look at how the network learns. Forward propagation consists of the step where the neural network performs a prediction vector based on an input feature vector x . The prediction vector will here be termed as h . The dot product of the weight matrix W (explained in the previous section) and the input vector x is then computed, which gives the vector z .

$$\begin{aligned}\vec{x}^T \cdot W &= \begin{pmatrix} x_1 & x_2 \end{pmatrix} \cdot \begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{pmatrix} \\ &= \begin{pmatrix} x_1 w_{11} + x_2 w_{21}, & x_1 w_{12} + x_2 w_{22}, & x_1 w_{13} + x_2 w_{23} \end{pmatrix} \\ &= \begin{pmatrix} z_1 & z_2 & z_3 \end{pmatrix} = \vec{z}, \vec{h} = \sigma(\vec{z})\end{aligned}$$

In order to obtain the prediction vector h , one of three non-linear activation functions is applied to the dot product vector z . The three activation functions can be either tanh, sigmoid or ReLu, which enables a non-linear mapping from z to h . This allows the network to understand non-linear data through the linear combination value of the weights and neurons and the non-linear activation function.

The last part of the learning process of the neural network, and the last step in forward propagation, consists of computing a loss function L . This phase consists of comparing the prediction vector y with the ground truth \hat{y} . The ground truth label represents the actual values, where the difference between the values of \hat{y} and the predicted values y are computed in order to measure the network's accuracy. Minimizing the loss function L enables the network to learn and adjust the weights accordingly, as the distance between y and \hat{y} is revealed to the network. In other words, the network is able to learn how the parameters θ contribute to the loss function and how it should change these parameters in order to decrease the distance between \hat{y} and y .

Two of the most common loss functions implemented in deep learning are the mean squared error loss and the cross-entropy loss.

Mean squared error loss:

$$L(\theta) = \frac{1}{N} \sum_{i=0}^N (y_i - \hat{y}_i)^2$$

Where,

y_i = entries in the prediction vector \vec{y}

\hat{y}_i = entries in the ground truth label $\hat{\vec{y}}$

Cross-Entropy Loss:

$$L(\theta) = -\sum_{i=0}^N \hat{y}_i \cdot \log(y_i)$$

Where,

y_i = entries in the prediction vector \vec{y}

\hat{y}_i = entries in the ground truth label $\hat{\vec{y}}$

[Oppermann, 2019].

Minimizing the loss function is solved mathematically by the method termed gradient descent. The method involves taking the derivative of the loss function where each gradient descent step enables the network to get closer in optimizing the weights, and ultimately make satisfying predictions. A simplified example would be a neural network with one input neuron, one output neuron and one connecting weight value, where the network first performs an inaccurate prediction based on a certain weight value. To improve the accuracy a loss function needs to be chosen (in this case a quadratic loss function) where the gradient descent method will update the weight with a more correct value.




Figure 2.4: Gradient Descent

The x-axis in figure 2.4 is the value of the weight, while the y-axis is the difference between the ground truth and the prediction, and essentially the network parameters (the value of the loss function). Using a quadratic loss function means that there is only a global minimum and not a local minimum. The global minimum represents the value of the weight that enables the network to make an accurate prediction, and is shown in the figure as the point where the loss function is equal to zero. In the fictive example shown in the figure, the initial prediction and weight value is relatively inaccurate. To improve the accuracy of the prediction, the next step involves updating the weight parameter until it reaches an optimized value for the weight. This is carried out by taking the derivative of the loss function with respect to the weight w .

$$\nabla_{w11} L(w11) = \frac{\partial L(w11)}{\partial w11}$$

The equation computes the tangent (slope) of the loss function for the point of where the initial weight lies. The figure shows that the tangent of the loss function would be positive, which would correspond to the network updating the weights to a higher loss value and making even more inaccurate predictions. However, this can be solved by multiplying the gradient by minus one to get the opposite direction of the gradient.

$$w_{11_{new}} = w_{11_{old}} - \epsilon \cdot \nabla_{w_{11}} L(w_{11})$$

The equation above shows a step of the gradient descent where the network updates the value of the weight to make a more accurate prediction. In the equation, epsilon is implemented as the network's learning rate and represents how quickly the parameters are updated. This method of updating the parameters continues until the loss function reaches a point of global minimum, which represents the optimal weight value for the network to perform accurate predictions [Oppermann, 2019].

2.5.4 LSTM Network

Unlike feed forward neural networks (FNNs) explained in the previous section, long short-term memory networks (LSTMs) is a type of recurrent neural network (RNN). Comparing RNNs with FNNs, RNNs resemble a more accurate representation of the workings of a biological neural network. The human brain has the ability to retain previous learnt information while receiving new information, which can either change a perceived outcome or it can remain unchanged. The same logic applies to RNNs, where loops within the networks are utilized to incorporate past information in the final output [Adusumilli, 2015].




Figure 2.5: Structure of a Simple RNN

Figure 2.5 illustrates a simple RNN where A represents a piece of neural network (repeating module), X_t is some input, h_t an output value, while the loop allows the network to pass information from each step to the next. The right side of the

figure illustrates a recurrent neural network if it had been unrolled, and demonstrates a chain-like characteristic useful for data consisting of sequences and lists [Olah, 2015].

A limitation with simple RNNs is their ability to retain previous learnt information over longer periods of time. Their short-term memory is connected to the vanishing gradient problem [Phi, 2015]. During backpropagation, these simple RNNs' gradients shrink as they back propagate through time and their values become extremely small. This applies normally to the earlier layers when the update values become so small that the layers stop learning. This can make the RNNs forget what they have previously seen. In relation to predicting stock returns, LSTM is a type of RNN which can cope with this problem since financial time series requires predictive models to train on large data sets, stretching over a long period of time.




Figure 2.6: Structure of a LSTM-Network (source: [Olah, 2015])

Figure 2.6 illustrates the recurrent system of an LSTM. The yellow boxes represent learned network layers, the pink circles are pointwise operations, while each line visualizes how the LSTM carries entire vectors. Merging lines depict concatenation and lines that fork is a representation of copied contents going to other locations [Olah, 2015].




Figure 2.7: The Repeating Module (source: [Olah, 2015])

One of the main differences between a simple RNN and an LSTM is the repeating module. A standard RNN only has a single neural network layer, while an LSTM has four interacting neural network layers. The horizontal line running through figure 2.7 illustrates the cell state C_t . This represents the flow of information through the network chain. In order for the LSTM to regulate the information being passed through, three different types of gates are used. The gates are made up of a neural network layer and a pointwise multiplication operation, and consist of a forget gate f_t , an input gate i_t , and an output gate o_t [Olah, 2015]. The forget gate is a sigmoid layer that decides which parts of the information to forget. f_t outputs a value between 1 and 0 for each number in the cell state C_{t-1} by evaluating h_{t-1} and x_t , where a value of 1 means keeping all of the information and a value of 0 means forgetting all of the information:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

The next two parts of the LSTM decide what new information to be stored in the cell state. The first part consists of an input gate where a sigmoid layer decides which values to update [Olah, 2015]. Then, a vector of new candidate values C_t are created by a tanh layer, which represents potential values to be added to the state:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Having determined the information to forget and keep from memory, how much information that will be added to the memory, and candidate values to be added, the next step involves updating the cell state:

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

The old state is multiplied with the forget gate layer f_t , while the new, scaled candidate values $i_t * \tilde{C}_t$ are added. The final step is for the LSTM to decide on the output. A sigmoid layer determines which parts of the cell state to output, before the cell state runs through a tanh layer in order to transform the values into values between -1 and 1 [Olah, 2015]:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

2.5.5 GRU

The Gated Recurrent Unit (GRU) is a variation of an RNN and an LSTM network. Compared to LSTM, the GRU is a more simple model where the forget and input gates are combined into an update gate, as well as having merged the cell state and hidden state [Olah, 2015].




Figure 2.8: Structure of a Gated Recurrent unit (source: [Olah, 2015])

In the update gate, x_t is plugged into the network unit and multiplied by the

weight W_z . h_{t-1} holds previous information and is also multiplied by its weight U_z . Further, a sigmoid activation transforms the result into a value between 0 and 1 [Kostadinov, 2017]:

$$z_t = \sigma(W_z x_t + U_z h_{t-1})$$

Similar to the LSTM, in this step the GRU determines which parts of the past information to pass along in the structure, where a value of 1 means copying all of the past information and a value of 0 means keeping none of the information. Further, the reset gate is used to determine which parts of the information to forget. The same logic of the update gate applies to the reset gate:

$$r_t = \sigma(W_r x_t + U_r h_{t-1})$$

The next step of the process is introducing a memory content, in order for the network to be able to utilize the reset gate to store relevant information from the past [Kostadinov, 2017]:

$$\tilde{h}_t = \tanh(W_{x_t} + r_t * U_{h_{t-1}})$$

The equation allows the model to determine what to remove from previous learned information, before applying the nonlinear activation function \tanh , to transform the results into values between -1 and 1. The last step of the GRU is using the update gate to decide what information to collect from the current memory content \tilde{h}_t and the previous steps h_{t-1} [Kostadinov, 2017]. In order to do so, the model calculates h_t , which is a vector with information for the current unit and with a function of passing the information down in the network structure:

$$h_t = z_t * h_{t-1} + (1 - z_t) * \tilde{h}_t$$

2.5.6 CNN

Convolutional neural networks (CNNs) is a type of feedforward neural network commonly used for image recognition and object classification. The structure of a CNN can be compared to the neurons in the human brain, where individual neurons respond to stimuli in several Receptive Fields (parts of the visual field) that overlap each other to cover the entire visual area [Saha, 2018].




Figure 2.9: Structure of a CNN (source: [Zakka et al., 2021])

The architecture of a CNN differs from standard neural networks in terms of making an assumption of input being images, which reduces the amount of parameters in the model. Figure 2.9 displays the difference between them where the neurons in a CNN are arranged in three dimensions, which consist of width, height and depth (illustrated in the third layer). The third dimension of depth refers only to an activation volume and not the depth of an entire neural network [Zakka et al., 2021]. In addition to having fully connected layers, as in standard neural networks, the CNN uses convolutional layers and pooling layers to reduce spatial size of the convolved feature.

In the convolution layer a kernel K performs a matrix multiplication operation of K and P (the part of the image the kernel hovers over) equal to the number of stride lengths. The process is repeated until the image is traversed [Saha, 2018]. A CNN can have multiple convolutional layers, where the first one uses the input image to extract low-level features, while the last ones enables the model to adapt to the high-level features. Further, pooling layers reduce the amount of computational power needed to process the data, by reducing the spatial size of the convolved feature and by reducing dimensionality. There are two types of pooling, namely max pooling and average pooling, which consist of the model returning either the maximum value from the portion of the image covered by the kernel, or the average of all the values [Saha, 2018]. This is useful for extracting dominant features, as well as suppressing noisy activations. The last steps of the CNN-architecture consist of flattening the final output, before feeding it into a neural network (fully connected layer) to perform the classification.

2.6 Backtesting

Developing new trading strategies can be a difficult task, and especially for investors to gain the confidence to implement the strategies in real-life trading. Backtesting, or systems testing, is an important method used by brokers and professional investors to test new strategies in a simulated environment. Backtesting reconstructs a trading scheme on historical data that would have occurred by using a particular trading strategy. An assumption within backtesting is that if a trading strategy has performed well in the past, it will also probably perform well in the future. Vice versa, if a trading strategy has performed poorly in the past, there is a good chance it will perform poorly in the future as well [Ni and Zhang, 2005]. To evaluate the performance of a trading strategy, several statistics and metrics can be used. This includes net profit or loss, volatility measures, average gain or loss, the percentage of capital exposed to the market, win-loss ratios, annualized return, and risk-adjusted return [Kuepper, 2021].

Backtesting trading strategies comes with several potential pitfalls that can cause a seemingly successful strategy to fail when applying it to real-life trading:

- **Market trend:** The time period and broad market trend are important factors to consider when conducting backtesting. If a strategy is backtested on a time period characterized by a consecutive bull market, the chances are the strategy will not perform well in a bear market. Backtesting should therefore be carried out on historical data with market conditions that include both a positive and negative market sentiment.
- **Universe:** Backtesting should also consider the universe the system is tested in. Developing a trading strategy targeted towards small cap stocks, and backtesting it in a universe of large cap stocks, will probably not translate into developing a successful strategy. It is therefore important to be aware of the limitations of a given strategy, while making sure the backtesting complies with these limitations.
- **Volatility:** Traders who have leveraged accounts should focus their attention on volatility. Leveraged investors are subjected to margin calls, which occurs when an investor's account value drops below the broker's threshold. It is recommended for these investors to include volatility measures where the backtesting results in low volatility to minimize risk.
- **Risk:** Risk: A common mistake when backtesting is only focusing on

annualized returns without taking into account the change in risk. Annualized returns is often used as a metric to compare the performance of a strategy against other suitable benchmark strategies. Risk-adjusted returns is an alternative metric which should be implemented in backtesting, in order to be certain that a potential trading strategy actually outperforms other benchmark strategies with an equal (or lower) level of risk.

- **Storytelling:** In situations where an investor has worked hard to develop a new trading strategy it is not uncommon to be biased. Since backtesting involves testing strategies on historical data (ex-post), it is possible that the investor justifies certain patterns and results with rationality that skews the results in the investor's favour. This could potentially lead to applying strategies that underperform against other more well-established trading strategies.
- **Transaction costs:** Simulating the transaction costs of each trade being made in backtesting is difficult because it would require the trades to be actually carried out under the
- **Transaction costs:** Simulating the transaction costs of each trade being made in backtesting is difficult because it would require the trades to be actually carried out in the real world. The bid-ask spread varies greatly between individual stocks, which can be hard to simulate correctly.
- **Over-optimization:** This is a condition where performance results are overfitted to the past. There is no guarantee the past will repeat itself, which means that fine-tuning parameters too much to the past can cause a strategy to fail in the present and future.

Investors should keep these potential pitfalls in mind when backtesting a new trading strategy. Even if these frequent backtesting errors are avoided, it does not mean a good performance backtest will result in successful performance in the future. There will always be a chance of selection bias and a chance of history not repeating itself.

2.7 Transaction Costs

As mentioned in the previous section, accurately simulating transaction costs is a common error when developing a trading strategy. To be able to accurately

simulate the transaction cost for each given trade, it would require traveling back in time and actually executing each trade. To best cope with this problem, an investor needs to be aware of the different elements of transaction costs and how these elements affect the overall trading performance. The elements consist of:

- **Explicit transaction costs:** Broker commission and securities lending.
- **Implicit transaction costs:** Arrival cost and implementation shortfall.

[J.P. Morgan, 2018]

Explicit transaction costs can be decomposed into broker commissions and securities lending. Broker commissions is a fee an investor is required to pay in order to have a brokerage company facilitate a transaction. These fees vary between brokers, but it can have a substantial impact on returns, especially when a strategy requires frequent trading. Securities lending is directed at short selling activities. Short selling is when an investor borrows another investor's shares and sells these shares in the market, where the purpose is to buy back the same amount of shares at a later point at a lower price, before delivering the stocks back to the investor. Short sellers pay a borrowing fee to the investor who agreed to lend out the shares.

Implicit transaction costs can be decomposed into arrival cost and implementation shortfall. Arrival cost is the difference between the arrival price and the execution price. The arrival price is the given price an asset is valued at the moment right before an order, while the execution price is the actual traded price [J.P. Morgan, 2018]. Implementation shortfall is the opportunity cost of placing an order below ask and when this order does not get executed. Placing an order below ask might lead to missing the opportunity of taking a position in a stock, which can affect portfolio composition and returns. This part of trading can be quite difficult, where an investor has to decide whether to avoid implementation shortfall all together by filling the current orders at ask, or placing an order below ask with hopes of getting a better entry price.

A common way of backtesting deep learning trading algorithms is only including direct transactions costs. This underestimates the total transaction costs of a trading strategy, and can result in over-optimistic returns and results from the backtesting. Previous academic papers have usually assumed direct transaction costs of five basis points [Fischer and Krauss, 2018], while research carried out

in 2009 on the Oslo Stock Exchange, showed arrival costs of 200 basis points by studying the period between 2000-2008 [Ødegaard, 2009]. Since implicit transaction costs can potentially have a substantial impact on returns, it is important to include these costs in backtesting as an investor would experience arrival costs in practice.

2.8 Market Capitalization

The market capitalization of a stock, also frequently mentioned as market cap, is the total market value of a company. It is a measure based on a company's worth in the open market and can be found by multiplying the number of outstanding shares with the last traded share price. How the open market determines a company's market cap is dependent on several factors. A common misconception is that it measures equity. The market evaluation of a company's total value is dependent on several factors. One of the most important factors is the company's ability to generate cash flow. This includes the ability of generating cash flow both in present time and in the future. The market expectation of a company's ability to generate future cash flow is one of the main reasons why a company sometimes can be either overvalued or undervalued. These situations of incorrect short-term evaluation create opportunities of possible profitable trades, which could potentially be exploited by deep learning algorithms.

2.8.1 Large Cap Stocks

Stocks reviewed as large cap companies are big, solid companies with a stable positive cash flow, often associated with low risk. These stocks are referred to as blue chips, which originates from casinos where blue chips are the most valuable. Blue chips are often associated with solid earnings, low level of debt ratio and frequent cash dividends. Large cap stocks are often preferred by investors seeking a low level of risk, high volume of traded shares, and expectations of small fluctuations (low volatility) in share price. The Oslo Stock Exchange (OSE) is a smaller stock exchange which means that the classification of large cap stocks differ from larger stock exchanges such as S&P 500, Dow Jones and Nasdaq in the US. These typically classify companies with a market capitalization of minimum 10 billion dollars as large cap stocks, while at the Oslo Stock Exchange this boundary is set to a market cap above 15 billion NOK (approximately 1.8 billion dollars).

2.8.2 Small Cap Stocks

Companies regarded as small cap stocks are small, young companies, where a stable positive cash flow is not expected until several years down the line. The share prices tend to be more volatile with less liquidity than mature companies. These types of companies are often expected to experience significant growth in future years, which makes them difficult to evaluate. Because a positive net cash flow is not expected before many years, investing in these companies brings increased risk. Small caps are therefore often subject to speculations, where the share price can fluctuate greatly, based on changes in the company outlook and in periods of economic slowdowns. The stock exchanges in the US regard companies with a market cap below 2 billion dollars as small cap stocks, while the Oslo Stock Exchange regards companies with a market cap below 1 billion NOK (approximately 120 million dollars) as small cap stocks.

Comparing historical returns favours investing in small cap stocks over large cap stocks. In the period from 1997 to 2012 the small cap index Russel 2000 had an annual return of 8.6%, while the S&P 500 had an annual return of 4.8% [Segal, 2020]. On the other hand, the small cap index had one-third higher volatility compared to the S&P 500. Comparing small cap funds with large cap funds in the period 2003-2013 showed a standard deviation (volatility measure) of 19.28 for small cap funds and 15.54 for large cap funds. A higher volatility translates into higher risk, which can make it difficult for investors, who have a limited number of stocks in their portfolio, to capitalize from investing in small cap stocks.

Even though small cap stocks seem to outperform large cap stocks, they have higher volatility that translates into higher risk. Small cap investing with a portfolio consisting of a limited number of stocks can make it difficult for investors to capitalize from this type of investment strategy.

Data

Methodology

The following chapter elaborates on the methods and techniques used to compare the returns of applying deep learning algorithms to small cap- and large cap stock predictions. It presents a variety of deep learning models used for predicting returns for both large cap and small cap stocks on the Oslo Stock Exchange. The aim of the research is to evaluate if the excess returns of investing in small caps can be capitalized on, through deep learning algorithms being able to detect patterns and relationships within historical data and by making accurate predictions of the probability of future share price development. An important part of a research process is choosing the appropriate method and research design to correctly analyze information and data on the topic, as well as for others to be able to assess the reliability and validity of the paper.

The methodology and techniques used in this paper is inspired by the work of Krauss et al. and Lund et al. [Fischer and Krauss, 2018] [Lund and Løvås, 2018]. This means that a similar approach and framework for developing the algorithms and assessing the models' predictive performance is applied to predicting small cap stock returns. In addition to exploring the models' applicability to small cap stock returns, the paper will add on previous work in terms of implementing a gated recurrent unit (GRU) and a convolutional neural network (CNN).

The methodology can be decomposed into five sections, where the first three presents how the networks are trained and tested, how the dependent variable is created, the different independent features that are included, and the networks' hyperparameters. Further, an explanation of how the models' predictive performance are evaluated is presented, before defining different trading strategies and their applicability to a real-life trading scheme.

4.1 Long Short-Term Memory Network

4.1.1 Training and Testing

The process of training the algorithms starts by defining a study period based on the collected data set, which contains data from January 1st 2005 to December 31st 2020. This involves splitting the data into a training set and a test set. The training set has an interval of 750 trading days, while the test set has an interval of 250 trading days. Given that stock markets are closed on the weekend and during holidays, 750 trading days translates into approximately 3 years while 250 trading days translates into 1 year. The training period is where the LSTM-network learns the patterns within the data and adjusts parameters to increase predictive accuracy, while the trading sets (test sets) are used to make out-of-sample predictions. The study periods are divided into a total of 23 unique periods, where each period is set up as rolling blocks of 1000 days. This means that the first period of trading (testing) starts 750 days into the data set, where these first days are lost to training. Further, each study period begins 250 days after the previous one, to ensure that the trading periods are non-overlapping, which means testing the models' predictions on the same data is avoided [Fischer and Krauss, 2018].

The training and testing sequences are applied to two different groups of stocks, where one is a group of small cap stocks that are constituents of the OSESX small cap index and the other is a group of large cap stocks that are constituents of the OSEBX benchmark index. For both groups, n_i denotes the number of stocks that are part of indices at the last day in study period i . Some of the stocks may not have a full data history for each training period, due to either delisting, being listed at a later time, or implemented in another index. If stocks do not exhibit share price data after a certain point in the trading period, they are included for trading up until this point.

4.1.2 Independent Variables

In order to improve model performance in terms of finding patterns and relationships within the data, several independent variables are used as input features for the models to train and test on. A large segment in the Norwegian economy is based on oil-related companies and represents a substantial part of the Oslo Stock Exchange's total market cap. Brent-oil price is therefore regarded as an important influence on the Oslo Stock Exchange and will be included as an input

feature. Other independent variables that represent macroeconomic factors are the USD/NOK exchange rate and government bonds, which are also included as input features. Additional input features are based on technical indicators and consist of the stocks' daily volume, and moving averages with intervals of 50 and 200 days. Lastly, a measure of the broad market volatility is included as an input feature through daily data on the CBOE Volatility Index (VIX), also known as the fear index. The VIX captures market sentiment by measuring the relative strength of near-term price changes of the S&P 500.

4.1.3 Target Variable

4.1.4 Hyperparameters

4.2 Gated Recurrent Unit Network

4.2.1 Training and Testing

4.2.2 Independent Variables

4.2.3 Target Variable

4.2.4 Hyperparameters

4.3 Convolutional Neural Network

4.3.1 Training and Testing

4.3.2 Independent Variables

4.3.3 Target Variable

4.3.4 Hyperparameters

4.4 Performance Metrics

4.5 Strategies

4.6 Predictive Models

4.6.1 Benchmark Algorithms

Random forest is a classification algorithm which consists of a large number of individual decision trees that works like an ensemble. Decision trees split the data into branches, or subsets, where certain learning conditions enable the algorithm to classify observations. Random forest operates like an ensemble since it uses multiple algorithms to improve the performance and accuracy of the predictions. In a binary case, each individual tree performs a class prediction, where the class with most predictions is selected as the model's prediction. Combining several decision trees reduces the probability of overfitting and situations

where the model finds optimal solutions that are local instead of global. Logistic regression is another algorithm used for classification problems. It resembles linear regression, but uses a more complex cost function known as the Sigmoid function to return a probability value [Pant, 2019].

4.6.2 Deep Learning Algorithms

The deep learning algorithms which will be used in the paper will consist of variations of neural networks that are based on both a feedforward architecture and recurrent architecture. Feedforward neural networks are regarded as simpler than recurrent neural networks since the information only moves forwards and not backwards through the network. In contrast, recurrent neural networks have connections between nodes characterized as a directed graph, which forms a cycle. The network is thus able to retain its state while processing the next sequence of inputs. Current suggestions for algorithms to be incorporated in the paper consist of simple recurrent neural networks (SRNN), gated recurrent units (GRU), and long short-term memory networks (LSTM). LSTM is a type of recurrent neural network and consists of cells, input and output gates, and a forget gate, which makes the algorithm well-suited for financial time series. GRU is also a form of recurrent neural network that possesses similar traits as LSTM. However, the algorithm does not have an output gate and therefore has fewer parameters than LSTM. Previous evidence suggests that the algorithm is suitable for smaller and less frequent data sets [Chung et al., 2014].

References

- [Adusumilli, 2015] Adusumilli (2015). Adusumilli, R. (2019). Machine Learning to Predict Stock Prices - Utilizing a Keras LSTM model to forecast stock trends, Retrieved from <https://towardsdatascience.com/predicting-stock-prices-using-a-keras-lstm-model-4225457f0233>). Accessed: 2021-12-05.
- [Banz, 1981] Banz, R. W. (1981). The relationship between return and market value of common stocks. *Journal of Financial Economics*, 9(1):3 – 18.
- [Bauman et al., 1998] Bauman, W., Conover, M., and Miller, R. (1998). Growth versus value and large-cap versus small-cap stocks in international markets. *Financial Analysts Journal - FINANC ANAL J*, 54:75–89.
- [Brownlee, 2020] Brownlee (2020). Brownlee, J. (2020). 4 Types of Classification Tasks in Machine Learning, Retrieved from (<https://machinelearningmastery.com/types-of-classification-in-machine-learning/>). Accessed: 2021-28-04.
- [Chen, 2020] Chen (2020). Chen, J. (2020). Risk-Free Rate Of Return, Retrieved from (<https://www.investopedia.com/terms/r/risk-free-rate.asp>). Accessed: 2021-10-03.
- [Chen, 2021] Chen (2021). Chen, J. (2021). Mean Reversion, Retrieved from (<https://www.investopedia.com/terms/m/meanreversion.asp>). Accessed: 2021-20-03.
- [Chung et al., 2014] Chung, J., Gulcehre, C., Cho, K., and Bengio, Y. (2014). Empirical evaluation of gated recurrent neural networks on sequence modeling. In *NIPS 2014 Workshop on Deep Learning, December 2014*.
- [Fama, 1970] Fama, E. F. (1970). Efficient capital markets: A review of theory and empirical work. *The Journal of Finance*, 25(2):383–417.
- [Fischer and Krauss, 2018] Fischer, T. and Krauss, C. (2018). Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research*, 270(2):654 – 669.

- [Foote, 2019] Foote (2019). Foote, Keith D. (2019). A Brief History of Machine Learning, Retrieved from (<https://www.dataversity.net/a-brief-history-of-machine-learning/>). Accessed: 2021-15-04.
- [Grossman and Stiglitz, 1980] Grossman, S. and Stiglitz, J. (1980). On the impossibility of informationally efficient markets. *American Economic Review*, 70:393–408.
- [Jensen, 1978] Jensen, M. C. (1978). Some anomalous evidence regarding market efficiency. *Journal of Financial Economics*, 6(2):95–101.
- [J.P. Morgan, 2018] J.P. Morgan (2018). J.P. Morgan Asset Management (2018). Transaction Costs Explained, Retrieved from (https://am.jpmorgan.com/blob-gim/1383537981326/83456/JPM50934_MiFID%20II%20Transaction%20Costs%20Guide_A5_FINAL.pdf). Accessed: 2021-07-05.
- [Kenton, 2021] Kenton (2021). Kenton, W. (2021). Beta, Retrieved from (<https://www.investopedia.com/terms/b/beta.asp>). Accessed: 2021-10-03.
- [Kostadinov, 2017] Kostadinov (2017). Kostadinov, S. (2017). Understanding GRU Networks, Retrieved from <https://towardsdatascience.com/understanding-gru-networks-2ef37df6c9be>). Accessed: 2021-13-05.
- [Krauss et al., 2017] Krauss, C., Do, X. A., and Huck, N. (2017). Deep neural networks, gradient-boosted trees, random forests: Statistical arbitrage on the sp 500. *European Journal of Operational Research*, 259(2):689 – 702.
- [Kryzanowski et al., 1993] Kryzanowski, L., Galler, M., and Wright, D. W. (1993). Using artificial neural networks to pick stocks. *Financial Analysts Journal*, 49(4):21.
- [Kuepper, 2021] Kuepper (2021). Kuepper, J. (2019). The Importance of Back-testing Trading Strategies, Retrieved from (<https://www.investopedia.com/articles/trading/05/030205.asp>). Accessed: 2021-06-05.
- [Lekovic, 2018] Lekovic, M. (2018). Evidence for and against the validity of efficient market hypothesis. *Economic Themes*, 56:369–387.
- [Linløkken, 2021] Linløkken (2021). Linløkken, G. (2021). Investtech - Forskningsresultater: Signalstatistikk, Retrieved from (<https://www.investtech.com/no/market.php?CountryID=1&p=>

staticPage&fn=helpItem&tbReport=h_SignalResultVideos). Accessed: 2021-15-03.

[Lund and Løvås, 2018] Lund and Løvås (2018). Lund, H. and Løvås, J. (2018). Employing Deep Learning for Stock Return Prediction on the Oslo Stock Exchange: A comprehensive test of deep learning models for predictive purposes and applicability to trading. Norwegian School of Economics, Retrieved from (<https://openaccess.nhh.no/nhh-xmlui/bitstream/handle/11250/2586263/masterthesis.PDF?sequence=1&isAllowed=y>). Accessed: 2021-01-01.

[Malkiel, 2003] Malkiel, B. G. (2003). The efficient market hypothesis and its critics. *Journal of Economic Perspectives*, 17(1):59–82.

[Metts et al., 2008] Metts, G., Slaubaugh, M., and Syed-Makander, A. (2008). Alpha investing: Exploring investment portfolio heuristics using real data introduction.

[Ni and Zhang, 2005] Ni, J. and Zhang, C. (2005). An efficient implementation of the backtesting of trading strategies. volume 3758, pages 126–131.

[Olah, 2015] Olah (2015). Olah, C. (2015). Understanding LSTM Networks – colah’s blog, Retrieved from <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>). Accessed: 2021-12-05.

[Olden, 2016] Olden (2016). Olden, M. (2016). Predicting Stocks with Machine Learning: Stacked Classifiers and other Learners Applied to the Oslo Stock Exchange. UIO, Department of Informatics, Retrieved from (<https://www.duo.uio.no/bitstream/handle/10852/51275/PredictingStocksWithMachineLearning.pdf?sequence=1&isAllowed=y>). Accessed: 2021-03-01.

[Oppermann, 2019] Oppermann (2019). Oppermann, A. (2019). What is Deep Learning and How does it work?, Retrieved from (<https://towardsdatascience.com/what-is-deep-learning-and-how-does-it-work-2ce44bb692ac>). Accessed: 2021-26-04.

[Panchal, 2018] Panchal (2018). Panchal, S. (2018). Artificial Neural Networks — Mapping the Human Brain, Retrieved from (<https://medium.com/predict/>

artificial-neural-networks-mapping-the-human-brain-2e0bd4a93160).
Accessed: 2021-27-04.

[Pant, 2019] Pant (2019). Pant, A. (2019). Introduction to Logistic Regression, Retrieved from (<https://towardsdatascience.com/introduction-to-logistic-regression-66248243c148>). Accessed: 2021-06-01.

[Phi, 2015] Phi (2015). Phi, M. (2018). Illustrated Guide to LSTM's and GRU's: A step by step explanation, Retrieved from <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb8>
Accessed: 2021-12-05.

[Poterba and Summers, 1988] Poterba, J. M. and Summers, L. H. (1988). Mean reversion in stock prices: Evidence and implications. *Journal of Financial Economics*, 22(1):27–59.

[Rompotis, 2019] Rompotis (2019). Rompotis, G. (2019). Large-cap vs small-cap portfolio performance: New empirical evidence from ETFs. *Review of Accounting & Finance*, 19(1), 71-94. doi:<http://dx.doi.org.ezproxy.library.bi.no/10.1108/RAF-03-2017-0056>, Retrieved from (<https://www.britannica.com/science/gauging-station>). Accessed: 2021-02-01.

[Saha, 2018] Saha (2018). Saha, S. (2018). A Comprehensive Guide to Convolutional Neural Networks — the ELI5 way, Retrieved from <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd>
Accessed: 2021-13-05.

[Segal, 2020] Segal (2020). Segal, T. (2020). Are Small Cap Companies Riskier Investments Than Large Cap Companies?, Retrieved from (<https://www.investopedia.com/ask/answers/022715/are-small-cap-companies-more-risky-investments-large-cap-companies.asp>). Accessed: 2021-08-05.

[White, 1988] White (1988). Economic prediction using neural networks: the case of ibm daily stock returns. In *IEEE 1988 International Conference on Neural Networks*, pages 451–458 vol.2.

[Zakka et al., 2021] Zakka et al. (2021). Zakka et al. (2021). CS231n Convolutional Neural Networks for Visual Recognition, Retrieved from <https://cs231n.github.io/convolutional-networks/>). Accessed: 2021-13-05.

[Ødegaard, 2009] Ødegaard (2009). Ødegaard, B.A. (2009). Hva koster det å handle aksjer på Oslo Børs?, Retrieved from (https://ba-odegaard.no/publications/2009_pof_hva_koster/handlekostnader_jan_2009.pdf). Accessed: 2021-07-05.